

Chan-Vese Image Segmentation

Shashank Sharma

December 20, 2023

1 Introduction

On a computer, images are represented using arrays of integers. Arrays for grayscale scale images have two dimensions, whereas for RGB images they have a 3rd dimension called *channel* as well, one for each of the colors Red, Green and Blue. The Image Segmentation problem is then to take such an array of integers and output a segmentation boundary that divides the image into the various objects it contains.

This is a very well researched problem and there exist various methods in literature for doing this. One of the most basic ones is thresholding - classifying pixels as *white* if they are above a threshold, and black instead. Canny edge detection is slightly more complicated - based on convolution against edge detection filters. Nowadays, deep neural network based techniques dominate the field.

In this report, we'll study the Chan-Vese model for image segmentation. It is a continuous model of the digital image, and the segmentation boundary is specified as the solution to a minimization problem. Level set methods can then be used to evolve an initial curve to a local minimum for the problem.

2 The Mumford Shah model

A given grayscale image is modeled as a function $f : \Omega \subset \mathbb{R}^2 \rightarrow [0, \infty]$. The Mumford and Shah model ^{MumfordShah} [9] approximates f by a piecewise-smooth function u as the solution to the minimization problem

$$\arg \min_{u, C} \mu \text{Length}(C) + \lambda \int_{\Omega} (f(x) - u(x))^2 dx + \int_{\Omega \setminus C} |\nabla u(x)|^2 dx, \quad (2.1)$$

where C is an edge set curve where u is allowed to be discontinuous. The first term ensures the regularity of C , the second term encourages u to be close to f , and the third term ensures that u is differentiable and has a small gradient on $\Omega \setminus C$.

This is a very natural way to pose segmentation, but algorithms for solving this minimization model tend to be very complicated and computationally expensive. Thus as a

simplification, Mumford and Shah also considered a piecewise constant formulation,

$$\arg \min_{u,C} \mu \text{Length}(C) + \lambda \int_{\Omega} (f(x) - u(x))^2 dx \quad (2.2)$$

where u is required to be constant on each connected component of $\Omega \setminus C$.

3 The Chan-Vese Model

The Chan-Vese model ^{[chan-veseorig](#)} [\[3\]](#) is a further simplification of the piecewise constant Mumford-Shah model. It restricts u to only have two values - c_1 inside C and c_2 outside C . It also adds an additional term penalizing the area inside C . Thus, the model becomes

$$\begin{aligned} \arg \min_{c_1, c_2, C} & \mu \text{Length}(C) + \nu \text{Area}(\text{inside}(C)) \\ & + \lambda_1 \int_{\text{inside}(C)} (f(x) - c_1)^2 dx + \lambda_2 \int_{\text{outside}(C)} (f(x) - c_2)^2 dx. \end{aligned} \quad (3.1)$$

$\mu, \nu, \lambda_1, \lambda_2$ are hyper-parameters to be determined by experimentation.

4 A semi-implicit gradient algorithm for solving the Chan-Vese minimization problem

To actually solve the minimization problem, it's convenient to represent C as the zero-crossing of a level-set function ϕ such that

$$C = \{x \in \Omega : \phi(x) = 0\} \quad (4.1)$$

and ϕ has different signs inside and outside C . Provided that ϕ is smooth enough and is indeed a distance function, the minimization problem can be rewritten as

$$\begin{aligned} \arg \min_{c_1, c_2, \phi} & \mu \int_{\Omega} \delta(\phi(x)) |\nabla \phi(x)| dx + \nu \int_{\Omega} H(\phi(x)) dx \\ & + \lambda_1 \int_{\Omega} (f(x) - c_1)^2 H(\phi(x)) dx + \lambda_2 \int_{\Omega} (f(x) - c_2)^2 (1 - H(\phi(x))) dx, \end{aligned} \quad (4.2)$$

where H is the heaviside function and δ the dirac delta distribution.

For a fixed ϕ , the optimal c_1 and c_2 are the region averages

$$c_1 = \frac{\int_{\Omega} f(x) H(\phi(x)) dx}{\int_{\Omega} H(\phi(x)) dx}, \quad (4.3)$$

$$c_2 = \frac{\int_{\Omega} f(x) (1 - H(\phi(x))) dx}{\int_{\Omega} (1 - H(\phi(x))) dx}. \quad (4.4)$$

For solving the problem numerically, H is approximated by the smooth function

$$H_\epsilon(t) = \frac{1}{2} \left(1 + \frac{2}{\pi} \arctan\left(\frac{t}{\epsilon}\right) \right) \quad (4.5)$$

and δ by its derivative

$$\delta_\epsilon(t) = \frac{d}{dt} H_\epsilon(t) = \frac{\epsilon}{\pi(\epsilon^2 + t^2)}. \quad (4.6)$$

This introduces another hyper-parameter ϵ .

For fixed c_1, c_2 , we can use the Euler Lagrange equations to come up with a gradient descent for optimizing ϕ

$$\frac{\partial \phi}{\partial t} = \delta_\epsilon(\phi) \left(\mu \operatorname{div} \left(\frac{\nabla \phi}{|\nabla \phi|} \right) - \nu - \lambda_1(f - c_1)^2 + \lambda_2(f - c_2)^2 \right) \text{ in } \Omega, \quad (4.7)$$

$$\frac{\delta_\epsilon(\phi)}{|\nabla \phi|} \frac{\partial \phi}{\partial \vec{n}} = 0 \text{ on } \partial\Omega, \quad (4.8)$$

where \vec{n} is the outward normal on the image boundary.

Now, we discretize the problem to solve it numerically. f is sampled on a regular grid $\Omega = \{0, \dots, M\} \prod \{0, \dots, M\}$. We discretize in space as (Method of Lines)

$$\frac{\partial \phi_{i,j}}{\partial t} = \delta_\epsilon(\phi_{i,j}) \left[\mu \left(\nabla_x^- \frac{\nabla_x^+ \phi_{i,j}}{\sqrt{\eta^2 + (\nabla_x^+ \phi_{i,j})^2 + (\nabla_y^0 \phi_{i,j})^2}} + \nabla_y^- \frac{\nabla_y^+ \phi_{i,j}}{\sqrt{\eta^2 + (\nabla_x^0 \phi_{i,j})^2 + (\nabla_y^+ \phi_{i,j})^2}} \right) \right. \quad (4.9)$$

$$\left. - \nu - \lambda_1(f_{i,j} - c_1)^2 + \lambda_2(f_{i,j} - c_2)^2 \right], \quad i, j = 1, \dots, M-1, \quad (4.10)$$

where ∇_x^- denotes backward difference, ∇_x^+ denotes forward difference and ∇_x^0 denotes central difference in the x dimension. Note that we have introduced another hyper-parameter η to prevent division by 0. The reason for mixing ∇_x^- and ∇_x^+ is to center the combined difference at (i, j) .

For convenience, we let

$$A_{i,j} = \frac{\mu}{\sqrt{\eta^2 + (\nabla_x^+ \phi_{i,j})^2 + (\nabla_y^0 \phi_{i,j})^2}}, B_{i,j} = \frac{\mu}{\sqrt{\eta^2 + (\nabla_x^0 \phi_{i,j})^2 + (\nabla_y^+ \phi_{i,j})^2}} \quad (4.11)$$

and use Forward Euler to discretize time. Then the fully discretized scheme looks like

$$\begin{aligned} \frac{\phi_{i,j}^{n+1} - \phi_{i,j}^n}{dt} = & \delta_\epsilon(\phi_{i,j}) [(A_{i,j}(\phi_{i+1,j} - \phi_{i,j}) - A_{i-1,j}(\phi_{i,j} - \phi_{i-1,j})) \\ & + (B_{i,j}(\phi_{i,j+1} - \phi_{i,j}) - B_{i,j-1}(\phi_{i,j} - \phi_{i,j-1})) \\ & - \nu \lambda_1(f_{i,j} - c_1)^2 + \lambda_2(f_{i,j} - c_2)^2] . \end{aligned} \quad (4.12)$$

fully discret

algo

Algorithm 1 A gradient descent based algorithm for Chan-Vese

```

Initialize  $\phi$ 
for  $n = 1, 2, \dots$  do
    Compute  $c_1, c_2$  as region averages
    Evolve  $\phi$  using one Forward-Euler time step
    if  $\frac{\|\phi^{n+1} - \phi^n\|_2}{\sqrt{|\Omega|}} < tol$  then
        break
    end if
end for

```

The boundary condition is enforced by creating ghost points so that

$$\phi_{-1,j} = \phi_{0,j}, \phi_{M,j} = \phi_{M-1,j}, \phi_{i,-1} = \phi_{i,0}, \phi_{i,M} = \phi_{i,M-1}. \quad (4.13)$$

We terminate when $\frac{\|\phi^{n+1} - \phi^n\|_2}{\sqrt{|\Omega|}}$ is less than a prescribed tolerance tol .

Algorithm [1](#) describes the full procedure.

To extend this algorithm to segment RGB images, we use the Chan-Sandberg-Vese model [\[4\]](#) instead. It just makes f, c_1, c_2 as d -dimensional, and the new minimization problem becomes

$$\begin{aligned} \arg \min_{c_1, c_2, C} & \mu \text{Length}(C) + \nu \text{Area}(\text{inside}(C)) \\ & + \lambda_1 \int_{\text{inside}(C)} \|f(x) - c_1\|_2^2 dx + \lambda_2 \int_{\text{outside}(C)} \|f(x) - c_2\|_2^2 dx. \end{aligned} \quad (4.14)$$

The overall algorithm remains the same.

5 Practical Implementation: The Code

The code written works for both grayscale and RGB images of any size.

First, we set the constants and create a function to display the images as shown in Figure [1](#). [fig:Code 1](#)

Next, we create functions to compute the dirac delta and heaviside functions, integrate and compute the average of a function using a specified density as shown in Figure [2](#). The function **extend** will be used to extend ϕ by one row / column on each side of the image, using 0-Neumann boundary conditions. The **restrict** function will be used to restrict any function back to the image domain. [fig:Code 2](#)

Next, we write functions for computing the forward, backward and central differences of a function in the row and column direction, as shown in Figure [3](#). The functions are designed so that the output is zero whenever calculating a difference involves indices outside the image bounds. So for example - the row backward difference along the first row of a function would be zero. [fig:Code 3](#)

```

1 pkg load image;
2
3 file = "apple.jpeg";
4
5 % Set constants
6 neta = 10^(-8);
7 mu = 0.2;
8 nu = 0;
9 lambda1 = 1; lamnda2 = 1;
10 dt = 0.5;
11 tol = 0.6*10^(-2);
12 global epsilon = 1;
13
14 %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
15
16 function display_image(f)
17     figure(1); clf;
18     if (length(size(f)) == 2)
19         colormap(gray);
20     endif
21     imagesc(f);
22     %imshow(f);
23     caxis([0 1])
24 endfunction
25

```

fig:Code 1

Figure 1: Setting constants and display-image function

```

25
26 function result = dirac_delta(A)
27     global epsilon;
28     result = epsilon ./ (pi * (epsilon^2 + A.^2));
29 endfunction
30
31 function result = heaviside(A)
32     global epsilon;
33     result = (1 + (2 * atan(A / epsilon) / pi)) / 2;
34 endfunction
35
36 function result = integrate(A)
37     result = sum(sum(A));
38 endfunction
39
40 function result = average(f, density)
41     result = integrate(f .* density) / integrate(density);
42 endfunction
43
44 %Extend f to outside the image boundary using 0 Neumann Boundary Condition
45 function result = extend(f)
46     A = [f(:,1, :), f, f(:,end, :)];
47     result = [A(1, :, :); A; A(end, :, :)];
48 endfunction
49
50 function result = restrict(f)
51     result = f(2:end-1, 2:end-1, :);
52 endfunction
53

```

fig:Code 2

Figure 2: Some basic functions defined

Next in Figure [fig:Code 4](#), we actually read the image from file and then rescale it so that pixel

```

53
54 %% Functions for various differences. Coded in such a way that the output
55 %% is 0 whenever a difference doesn't make sense
56
57 function result = row_forward_diff(f)
58     result = [f(2:end, :, :); f(end, :, :)] - f;
59 endfunction
60
61 function result = row_backward_diff(f)
62     result = f - [f(1, :, :); f(1:end-1, :, :)];
63 endfunction
64
65 function result = row_central_diff(f)
66     result = ([f(2:end, :, :); f(end-1, :, :)] - [f(2, :, :); f(1:end-1, :, :)]) / 2;
67 endfunction
68
69 function result = col_forward_diff(f)
70     result = [f(:, 2:end, :), f(:, end, :)] - f;
71 endfunction
72
73 function result = col_backward_diff(f)
74     result = f - [f(:, 1, :), f(:, 1:end-1, :)];
75 endfunction
76
77 function result = col_central_diff(f)
78     result = ([f(:, 2:end, :), f(:, end-1, :)] - [f(:, 2, :), f(:, 1:end-1, :)]) / 2;
79 endfunction
80

```

fig:Code 3

Figure 3: Functions for calculating various differences

```

82
83 %% Read an image from a file
84 f = imread(file);
85 info = imfinfo(file);
86
87 % convert image to double and scale to [0,1]
88 f = double(f);
89 f = f / 2^(info.BitDepth);
90
91 [n1,n2, n3] = size(f)
92 LargestPixel = norm(reshape(f, n1*n2*n3, 1), inf)
93
94 %%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%%
95
96 % Initialize phi
97 x = 1:n2; y = (1:n1)';
98 [xx, yy] = meshgrid(x, y);
99 phi = sin(pi * 10* x / n2) .* sin(pi * 10* y / n1);
100
101 display_image(f);

```

fig:Code 4

Figure 4: Reading the image from file and initializing ϕ

values lie between 0 and 1. ϕ is initialized using a scaled version of $\sin(\pi x)\sin(\pi y)$ as suggested in [3] because it leads to fast convergence.

Finally in Figure 5, we come to the Forward Euler time stepping. ϕ is first extended by one unit on each side of the image, and all the computations are carried out with this extended version to ϕ , so that boundary conditions are respected. The computations

```

102
103 step = 1;
104 while (1)
105
106     phi_extended = extend(phi); % Extend using 0 neumann bdd condns
107
108     phi_i_plus = row_forward_diff(phi_extended);
109     phi_i_minus = row_backward_diff(phi_extended);
110     phi_i_zero = row_central_diff(phi_extended);
111
112     phi_j_plus = col_forward_diff(phi_extended);
113     phi_j_minus = col_backward_diff(phi_extended);
114     phi_j_zero = col_central_diff(phi_extended);
115
116     A = mu ./ ((neta^2 + phi_i_plus.^2 + phi_j_zero.^2).^0.5);
117     B = mu ./ ((neta^2 + phi_i_zero.^2 + phi_j_plus.^2).^0.5);
118
119     d = heaviside(phi);
120     c1 = average(f, d);
121     c2 = average(f, 1 - d);
122
123     % Do all difference computations using phi_extended to enforce 0 Neumann bdd condns
124     Dphi = restrict(row_backward_diff(A .* phi_i_plus) + col_backward_diff(B .* phi_j_plus));
125     Dphi = Dphi - nu - lambda1 * sum((f - c1).^2, 3) + lambda2 * sum((f - c2).^2, 3);
126     Dphi = Dphi .* dirac_delta(phi);
127
128     phi_new = phi + Dphi * dt;
129
130     if (rem(step, 100) == 0)
131         step
132         display_image(f);
133         hold on;
134         contour(x, y, phi_new, [0,0], 'r');
135         drawnow
136         pause(0.001);
137         if (norm(reshape(phi_new - phi, n1*n2, 1), 2) / sqrt(n1*n2) < tol)
138             printf("Converged\n");
139             break;
140         endif
141     endif
142
143     phi = phi_new;
144     step += 1;
145
146 endwhile
147

```

fig:Code 5

Figure 5: Forward Euler time stepping

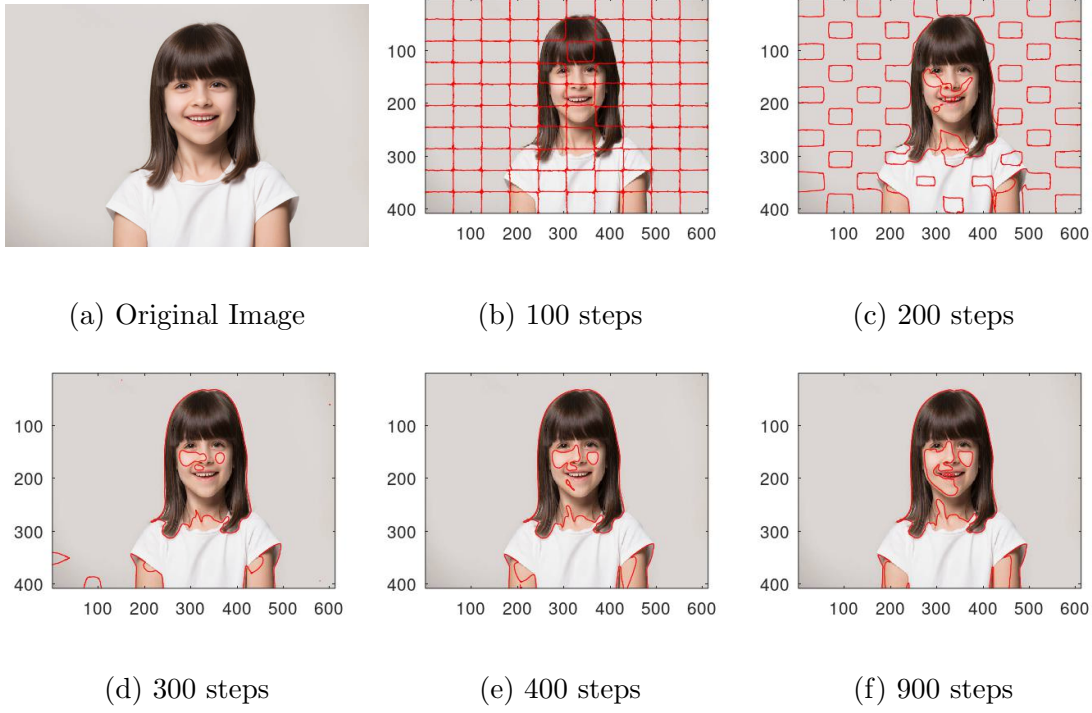
carried out are just a fully vectorized version of those described in [\(4.12\)](#). At every 100 steps, we display the contour on the image and also check to see if we satisfy the convergence criteria.

6 Examples

In all the examples, we use the parameters given in [Figure 1](#) unless stated otherwise. It's also important to note that the algorithm's runtime is very sensitive to the chosen values

Figure 6: Chan-Vese Segmentation on an RGB image

girl orig



of dt , initial condition of ϕ and the image size itself.

6.1 Effect of Noise

Figure [6](#) demonstrates the algorithm run on an RGB image. It takes 900 steps to converge.

Now, we add a random amount of noise to the image with strengths of 0.2, 0.5 and 1 respectively, and see how the algorithm performs. Figure [7](#) shows that the number of steps needed for convergence and the predicted contour remain largely unaffected, thus showing that the algorithm is fairly robust to noise.

6.2 Effect of μ


Figure [8](#) shows the effect of varying the length parameter μ . We see that as μ increases, the contour output becomes smoother as expected. However, the number of steps required by the algorithm to converge also increase considerably.

6.3 Effect of ν

Figure [9](#) shows the effect of varying the area parameter ν . For a negative ν , a bigger area inside the contour is rewarded, whereas for a positive ν it's penalized. We see that for

girl noisy




Figure 7: The algorithm's performance on noisy images



(a) Noise = 0.2

(b) Noise = 0.5

(c) Noise = 1




(d) 800 steps, Noise = 0.2

(e) 800 steps, Noise = 0.5

(f) 900 steps, Noise = 1

Figure 8: The effect of μ

galaxy



(a) Original Image

(b) $\mu=0.05$, 400 steps

(c) $\mu=0.2$, 700 steps

(d) $\mu=2$, 4200 steps

noisy apple




Figure 9: The effect of ν (a) Original (b) $\nu=-0.04$, 300 steps (c) $\nu=0$, 300 steps (d) $\nu=0.04$, 300 steps

Figure 10: Chan-Vese vs Canny

canny



(a) Original (b) 4800 steps (c) Canny thresh=0.3 (d) Canny thresh=0.5



(e) Original (f) 1000 steps (g) Canny thresh=0.2 (h) Canny thresh=0.5

$\nu = -0.04$, the contour almost covers the entire image to maximize the area reward. For $\nu = 0.04$, we see an additional contour boundary inside the apple that decreases the area.

6.4 Comparison against Canny

We compare the method against Canny-Edge detection on grayscale images, using Octave's *edge* function from the *image* package. We note that Canny Edge detection is always much faster than Chan-Vese, with Chan-Vese requiring several minutes of computation compared to Canny outputting in less than a second each time. Figure 10 shows the results.

7 Improvement Directions

One disadvantage of Chan-Vese is that it only segments the image into two phases. An extension of Chan-Vese for nested segmentation is ~~multilayer~~ ^{multilayer} [5] and multiphase segmentation is ~~multiphase1, multiphase2~~ ^{multiphase1, multiphase2} [10, 8]. Note that by the 3 color theorem, we only need at most 3 phases in general to segment any image.

Another disadvantage is that the gradient descent based algorithm above runs slowly. Other approaches include ~~topder~~ ^{topder} [7] based on the topological derivative, ~~multigrid~~ ^{multigrid} [1] based on the multigrid method and ~~graphcut1, graphcut2~~ ^{graphcut1, graphcut2} [6, 2] based on graph cuts.

References

- [1] Multigrid method for the chan-vese model in variational segmentation. *Communications in Computational Physics*, 4(2):294–316, 2008.
- [2] Egil Bae and Xue-Cheng Tai. *Efficient Global Minimization for the Multiphase Chan-Vese Model of Image Segmentation*, page 28–41. Springer Berlin Heidelberg, 2009.
- [3] Tony Chan and Luminita Vese. *An Active Contour Model without Edges*, page 141–151. Springer Berlin Heidelberg, 1999.
- [4] Tony F. Chan, B.Yezrielev Sandberg, and Luminita A. Vese. Active contours without edges for vector-valued images. *Journal of Visual Communication and Image Representation*, 11(2):130–141, June 2000.
- [5] Ginmo Chung and Luminita A. Vese. Energy minimization based segmentation and denoising using a multilayer level set approach. In Anand Rangarajan, Baba Vemuri, and Alan L. Yuille, editors, *Energy Minimization Methods in Computer Vision and Pattern Recognition*, pages 439–455, Berlin, Heidelberg, 2005. Springer Berlin Heidelberg.
- [6] Noha El Zehiry, Steve Xu, Prasanna Sahoo, and Adel Elmaghraby. Graph cut optimization for the mumford-shah model. In *The Seventh IASTED International Conference on Visualization, Imaging and Image Processing*, VIIP '07, page 182–187, USA, 2007. ACTA Press.
- [7] Lin He and Stanley Osher. *Solving the Chan-Vese Model by a Multiphase Level Set Algorithm Based on the Topological Derivative*, page 777–788. Springer Berlin Heidelberg.
- [8] Matthew Keegan, Berta Sandberg, and Tony Chan. A multiphase logic framework for multichannel image segmentation. *Inverse Problems and Imaging*, 1, 02 2012.
- [9] David Mumford and Jayant Shah. Optimal approximations by piecewise smooth functions and associated variational problems. *Communications on Pure and Applied Mathematics*, 42(5):577–685, July 1989.

- [10] Luminita A. Vese and Tony F. Chan. *International Journal of Computer Vision*, 50(3):271–293, 2002.

Shashank Sharma
Department of Mathematics, University of British Columbia, Canada.
sharm39@math.ubc.ca